

1 Files and Systems

1.1 How to Virtualize Persistent Storage

1.2 File Systems

- Provides long term storage
- Roles
 - Implements an abstraction (files) for secondary storage
 - Organize files logically (directories)
 - Permit sharing of data between processes, people and machine
 - Protect data from unwanted access (security)
- Requirements
 - Store a large amount of information
 - Information must survive after termination of using it
 - Multiple process must be able to access it concurrently
- Has two views
 1. User View
 - Convenient organization of information
 2. OS View
 - Managing physical storage
 - Enforcing access restrictions

1.3 File Operations

- Creation
- Writing
- Reading
- Repositioning within a file
- Deleting a file
- Truncation and appending
 - File attributes are kept
 - May erase/add parts or entire contents of a file

1.4 File Access Methods

- Sequential Access
 - Read bytes one at a time, in order
- Direct Access
 - Random access given block/byte number

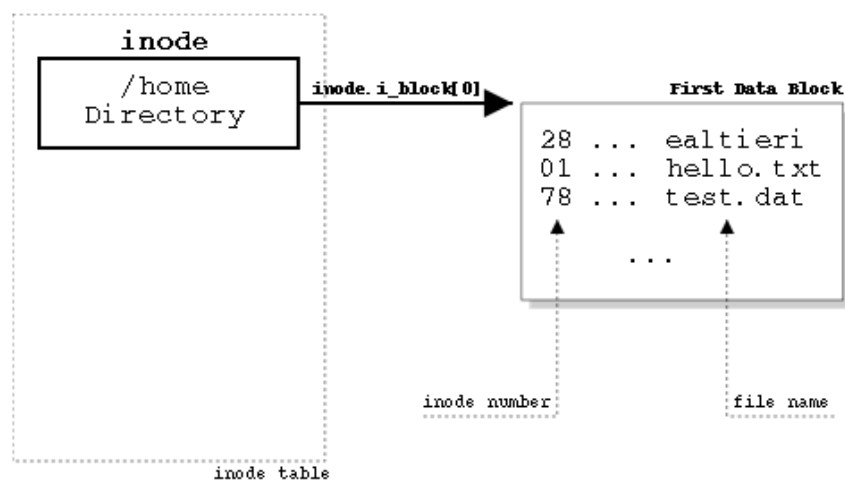
1.5 Directories

- Purpose
 - Provides logical structure to file system
 - * Users → A way to organize file
 - * File system → Convenient naming interface (programmers don't get lost)

1.6 Directory Structure

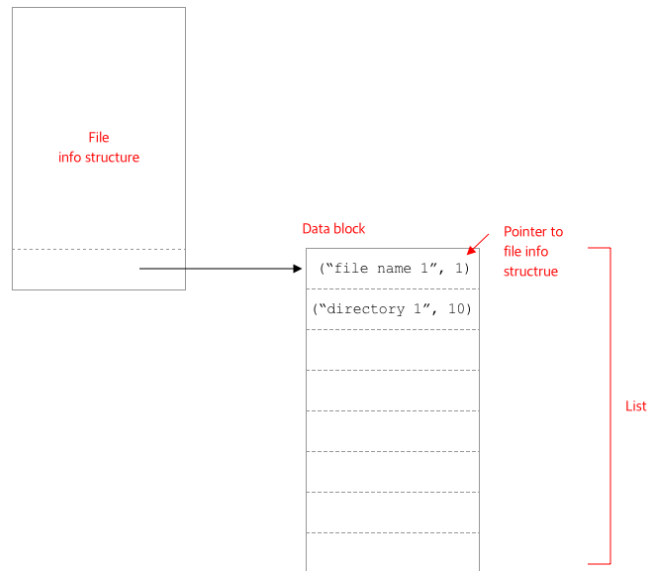
- Is a list of entries - names and associated metadata
 - List usually unordered
 - Metadata is like a label to a box
 - e.g size, protection, location, i-node number

Example



1.7 Directory Implementations

- Acyclic graph directories
- List
 - Simple list of file names and pointers to file info structure



- Hash table
 - Creates a list of file info structure
 - Hash file name to get a pointer to the file info structure

Example

```
directory["file-name"]
```

- Takes more space (hash table)

1.8 Hard Links

- Is a directory entry that associates a name with a file on a file system.
- Can't create new hard links to directories
 - May create cycles (Not good)
 - Breaks directory tree (Tree → no cycle!!)
 - Breaks unambiguity of parent directories
 - They multiply files
- Can't create hard links across partitions
- Removing a hard link only removes the file if it is the last link to file

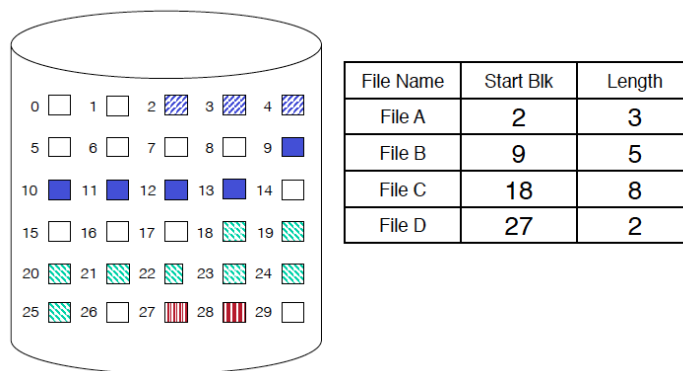
1.9 Symbolic Link

- Is a directory entry that associates a name with another directory entry on a file system.
- Can create a link to directory
- Can create a link across partitions
- File system needs to look up a link to follow it
- Removing a file → Dangling link

1.10 Associating Data Blocks to Files

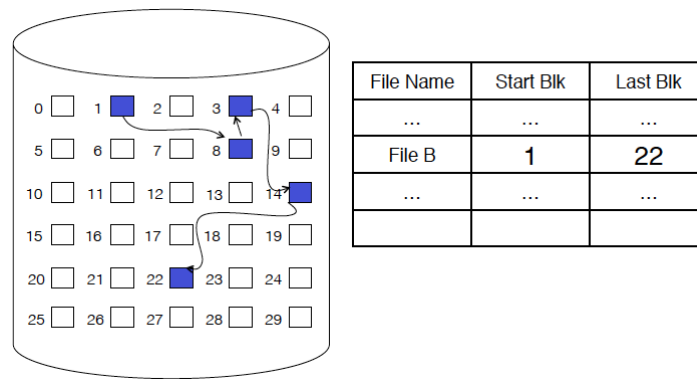
1.11 Disk Layout Strategies

- Contiguous Allocation




- Each file is given one (or more) set of contiguous set of blocks
- In info structure, stores
 - * Starting Block
 - * Length
- Pros and Cons
 - * Fast, simplifies directory access
 - * Allows indexing
 - * Inflexible
 - Need to know size of data blocks beforehand (when only one extent is allowed)
 - * Causes external fragmentation
- Example
 - * Microsoft NTFS
 - * Linux's ext4

- Linked Allocation



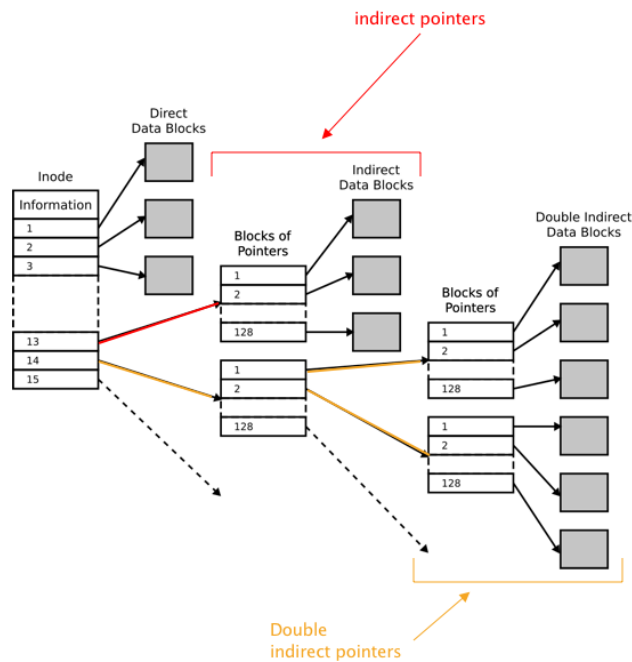
- Each block in file contains a pointer to next block
- In info structure, stores
 - * Starting Block
 - * Length
- Pros and Cons
 - * Good for sequential access
 - * Bad for others
- Example
 - * Microsoft FAT

- Indexed Allocation



- All file metadata is stored in an inode
- Each inode contains 15 pointers

- * 12 direct pointers
- * 1 single indirect pointer
 - Address of block containing address of data blocks
- * 1 double indirect pointer
- * 1 triple indirect pointer



- Pros and Cons
 - * No external fragmentation
 - * Handles random access better
 - * Files can be easily grown
 - * Requires many intermediary files
- Example
 - * Linux's ext2