

1 File System (cont')

1.1 Disk Layout Strategies

- How do you find all of the blocks for a file
 1. Contiguous Allocation
 - Using starting block + length (extent)
 - * All blocks of file are located together on disk
 2. Linked Allocation
 - Using Pointer
 - * Each block points to the next
 - * Directory entry points to the first
 3. Indexed-based allocation
 - Using pointer
 - * An "index" contains pointers to many other blocks
 - * May require multiple, linked index blocks

1.2 Data Block Allocation: Pros and cons

- Contiguous based allocation
 - Pros
 - * Fast sequential access
 - * Fast allocation
 - * Fast deallocation
 - * Small amount of metadata
 - Cons
- Question** What does it mean when moving whole files around?
 - * External fragmentation
 - * Needs compaction
 - * Inflexible
 - * Need to move whole files around
- Linked-based allocation
 - Pros
 - * Easy sequential access
 - * Disk blocks can be anywhere
 - * No external fragmentation
 - Cons

- * Direct access is expensive
- * If a data block is corrupted, could lose the rest of file
- Index-based allocation
 - Pros
 - * Handles direct access well (say block # 15 on file "I love corgi.txt")
 - * Disk blocks can be anywhere
 - * No external fragmentation
 - Data is placed randomly
 - Cons
 - * Limits file size ($\sim 4\text{TB}$)
 - * Cost of access bytes near the end of large files grows (i.e Triple indirect pointers)
 - pointer (in inode) \rightarrow data block of addresses (triple indirect pointers) \rightarrow data block of addresses (double indirect pointers) \rightarrow data block of addresses (single indirect pointers) \rightarrow data block

2 Implementation of Very Simple File System

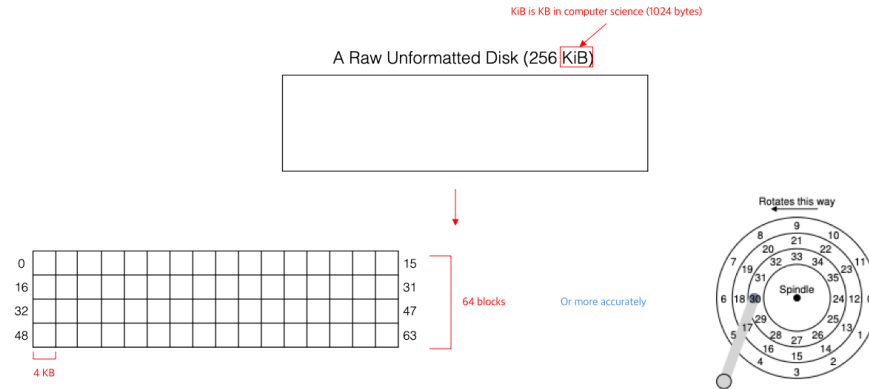
2.1 Existing File System

2.2 The main idea

- Goal: create a file system for an unformatted disk
- Key question: What does it mean to format a hard drive?
 - create some structure so that data will be easy to find and organize
- Other key questions:
 - Where do we store file data and metadata structures (answer: inode)
 - How do we keep track of data allocations
 - How do we locate file data and metadata
 - What are the limitations (max file size, etc.)?

2.3 File System Implementation

- Key Question: How does file system use the disk to store files?
 - By defining a block size (4 KB)
 - By allocating disk space in granularity of blocks



2.4 Superblock

- Key question: What do we need to know to connect the disk to a computer?
 - Using superblock!!
- Determines the location of root directory
- Is duplicated across disk for reliability (in case of corruption)
- Is at well known block
 - Is always read first (on boot) before attaching to file system
 - * Is to know where free map, and inodes are
 - * Is to know what kind of file system is being stored
 - * Is to know where the inode table begins
 - * The number of inodes and data blocks in a particular file system
 - * Other parameters

2.5 Freemap


- Key question: How do we keep track of data allocations? That is, how do we know which blocks are free and which blocks are in use?
 - Free map!!
- Determines which blocks are free
 - Usually bitmap, 1 bit per block on disk
 - Stored on disk, cached in memory for performance
 - * when read in memory, entries stored in L1 or L2 cache, so entries can be re-retrived quickly
- How many blocks can 4 Kib bitmap track of?
 - $4 \text{ KiB} = 32 \text{ Ki bits} \Rightarrow 32\text{Ki blocks}$

2.6 Data blocks

- Is used to store files and directories
- Is region used to store user data

2.7 Data Region

- Is where data blocks are located



2.8 Metadata: inode table

- Lets FS to keep track of information about each file
- **inode** has size of 128 bytes in VFSF
- How many files can VSFS hold at most?
 - $(5 \cdot 4\text{KiB})/128\text{B} = 160$ files

2.9 Allocation Structure

- keeps track of which blocks are being used and which ones are free
- data structure
 - Is called **bitmap**
 - 0 - free
 - 1 - not free
- Two types
 1. Inode bitmap (for inode)
 2. Data bitmap (for data blocks)

2.10 Indirection

- block of addresses - 4 bytes per pointer
- single indirect block
 - can fit 1024 data blocks
- double indirect block
 - can fit $1024 * 1024$ blocks
- triple indirect block
 - can fit $1024 * 1024 * 1024$ blocks
- How big can a file be?
 - $4\text{KiB} \cdot (12 + 1024 + 1024^2 + 1024^3) \approx 4\text{TiB}$

2.11 Why an imbalanced tree?

- Designed based on evidence
 - Because most files are small (2KiB)
 - Files are usually accessed sequentially
 - Directories are typically small (20 or fewer entries)

2.12 Another Approach: extent-based

- An extent == a disk pointer plus a length (in # of blocks)
- Instead of a pointer to every block of file, a pointer to every few blocks needed

Example

ext4, HFS+, NTFS, XFS

2.13 Yet another approach: Link-Based

- Uses in memory File Allocation Table, indexed by address of data block
 - Faster in finding a block
- Works poorly if we want to access the last block of a big file

Example

Microsoft's FAT file system

2.14 Summary

- Inodes
 - Data structure representing a FS object (file, dir, etc)
 - Attributes + disk location
 - No file name, just metadata
- Directory
 - List of (name, inode) mappings
 - Each directory entry: a file, other directory, link, itself (.), parent(..), etc