

1 Limited Direct Execution

Vocabulary

1. Time Sharing

- Is a mechanism used by an OS to share a resource
- Allows an entity to use the resource for a little while, and then a little while by another, and so forth

Example

CPU

2. Limited Direct Execution

- Is synonymous to **baby proofing**
- **Limited** Means there will be a limit to what a processor can and cannot do
- **Direct Execution** Means that the processor will run directly on the CPU

3. User Mode

- Is a processor mode where code that runs is restricted in what it can do

4. Kernel Mode

- Is a processor mode where code that runs can do what it likes, including privileged operations

Example

Privileged operations include

1. I/O requests
2. Executing all types of restricted instructions

5. System Call

- Is a programmatic way in which a computer program requests a privileged service from the kernel of the operating system

6. Trap

- Is a type of synchronous interrupt caused by an exceptional condition that
- Exceptional condition include:

- Breakpoint
- Division by zero
- Invalid memory access
- System Call
- Usually results in a processor switching to kernel mode

7. Return-from-Trap

- Is an instruction that
 - Restores saved registers from kernel stack
 - Switches the processor back to **user mode**

8. Trap Table

Question What is the exact definition of a trap table? OSTEP glosses over it :(

- Is synonymous to 대응 메뉴얼
- Is a list of trap handlers where each is associated with a specific trap

9. Trap Handlers

- Is the code that will run when the trap is triggered.

10. System-call Number

- Is an ID assigned to each system call

Linux System Call Table

The following table lists the system calls for the Linux 2.2 kernel. It could also be thought of as an API for the interface between user space and kernel space. My motivation for making this table was to make programming in assembly language easier when using only system calls and not the C library (for more information on this topic, go to <http://www.linuxassembly.org>). On the left are the numbers of the system calls. This number will be put in register %eax. On the right of the table are the types of values to be put into the remaining registers before calling the software interrupt 'int 0x80'. After each syscall, an integer is returned in %eax.

For convenience, the kernel source file where each system call is located is linked to in the column labelled "Source". In order to use the hyperlinks, you must first copy this page to your own machine because the links take you directly to the source code on your system. You must have the kernel source installed (or linked from) under '/usr/src/linux' for this to work.

System Call
Number

| %eax | Name | Source | %ebx | %ecx | %edx | %esi | %edi |
|------|-------------|--|----------------|-----------------------------------|------------------------|------|------|
| 1 | sys_exit | kernel/exit.c | int | - | - | - | - |
| 2 | sys_fork | arch/i386/kernel/process.c | struct pt_regs | - | - | - | - |
| 3 | sys_read | fs/read_write.c | unsigned int | char * | size_t | - | - |
| 4 | sys_write | fs/read_write.c | unsigned int | const char * | size_t | - | - |
| 5 | sys_open | fs/open.c | const char * | int | - | - | - |
| 6 | sys_close | fs/open.c | unsigned int | - | - | - | - |
| 7 | sys_waitpid | kernel/exit.c | pid_t | unsigned int * | int | - | - |
| 8 | sys_creat | fs/open.c | const char * | int | - | - | - |
| 9 | sys_link | fs/namei.c | const char * | const char * | - | - | - |
| 10 | sys_unlink | fs/namei.c | const char * | - | - | - | - |
| 11 | sys_execve | arch/i386/kernel/process.c | struct pt_regs | - | - | - | - |
| 12 | sys_chdir | fs/open.c | const char * | - | - | - | - |
| 13 | sys_time | kernel/time.c | int * | - | - | - | - |
| 14 | sys_mknod | fs/namei.c | const char * | int | dev_t | - | - |
| 15 | sys_chmod | fs/open.c | const char * | mode_t | - | - | - |
| 16 | sys_lchown | fs/open.c | const char * | uid_t | gid_t | - | - |
| 18 | sys_stat | fs/stat.c | char * | old_kernel_stat * | - | - | - |

11. Timer Interrupt

- Is a type of interrupt generated by an internal clock instead of an external event (e.g I/O or system call)

12. Interrupt Handler


- Is a special block of code associated with a specific interrupt condition

13. Disable Interrupts

- Ensures that when one interrupt is being handled, no other is delivered to CPU

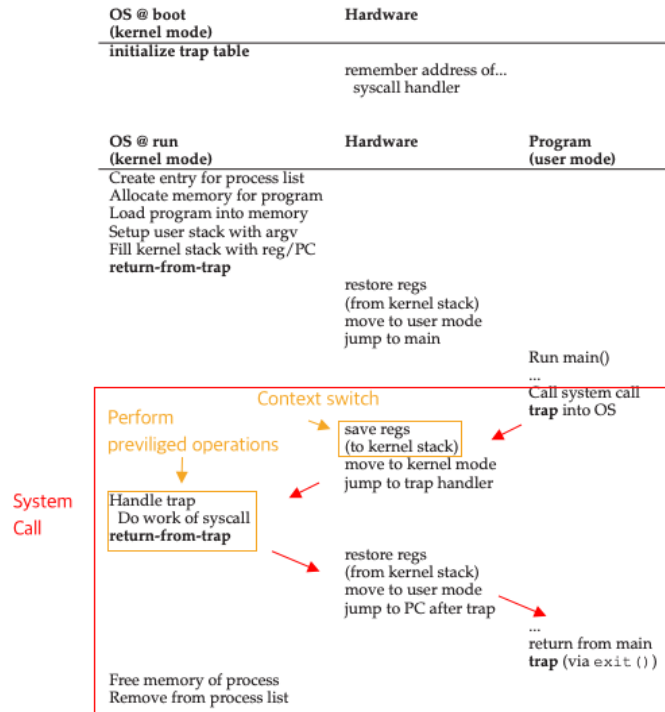
1.1 Direct Execution

- Means just run program directly without limits
- Has advantage of being fast



1.2 Problem #1: Restricted Operations

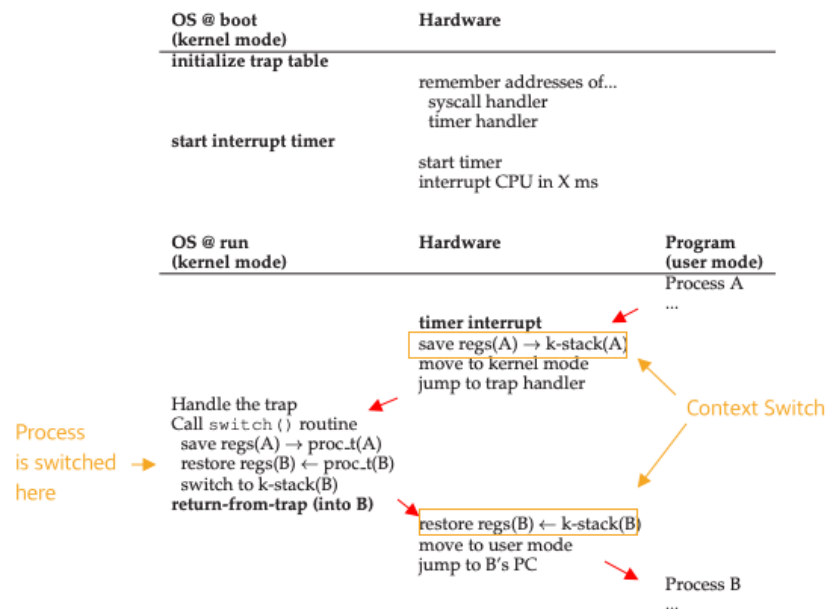
- **Question:** How can the OS make sure a program doesn't do anything that we don't want to do while running it efficiently?
- Solution
 - **user mode**
 - **kernel mode**
- **Question # 2:** What should a user process do when it wishes to perform some kind of privileged operation?
- Solution
 - **system call**
 - * Exact **system call number** is placed in a register or specified location on the stack
 - * OS, when handling the **system call**, examines the number, ensure its valid, and execute corresponding code



1.3 Problem #2: Switching Between Processes

- **Question:** When we are running a process, how does the operating system stop it from running and switch to another process, thus implementing **time sharing** mechanism to virtualize CPU?
- OS doesn't have control when process is running on CPU. So how can this be done?
- **Solution #1**
 - Wait for **System call** (Cooperative approach)
 - * Used in early days
 - * OS regains control from CPU when `yield()` system call is invoked
 - * Control also regained when other types of traps are raised
 - * Infinite loop → System call never invoked → Not good
- **Solution #2**
 - The OS takes control (A non-cooperative approach)
 - * Used today
 - * Uses **timer interrupt**
 - Interrupt every x milliseconds by a programmed timer device
 - Interrupt raised → interrupt handler in OS runs → OS regains control
 - Allows OS to stop current task and start a different one

- Saving and restoring context
 - Applies to both cooperative and non-cooperative approach
 - Steps
 - * **Timer interrupt** or (yield()) Is invoked
 - * **Scheduler** decides whether to continue running the current process or switch to a new one
 - Switch → **context switch!!**
 - * **Context switch** saves the following into current processe's kernal stack
 - General purpose registers
 - PC
 - Kernel stack pointer
 - * Restore registers and kernel stack of soon-to-be-executing process
 - * Call **Return-from-trap** instruction



1.4 Concurrency

- **Question:** How to handle an interrupt when one is being handled and another one happens?
- **Solution # 1:**
 - **Disable interrupt**
 - * Is a way to achieve mutual exclusion (Only one interrupt can be processed in critical section)

- * **Disable interrupt** during an interrupt processing so no other interrupts can enter
- * Enable interrupt once it leaves critical section